

Precision, Recall and F1 Score

```
from sklearn.metrics import recall_score, precision_score, f1_score
```

These three metrics are used to evaluate the performance of classification models, especially in scenarios **where the class distribution is imbalanced**.

Precision

- **Definition:** The proportion of **predicted positives that are actually positive**.
- Precision answers the question, *"Of all the instances that were predicted as positive, how many were actually positive?"*
- **Formula:** $\text{Precision} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Positives (FP)}}$

$$\text{Precision} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Positives (FP)}}$$

- **When to Use:** Use precision **when the cost of false positives is high**.

Example:

- In email spam detection, high precision means that most emails marked as spam really are spam.
 - You don't want to label legitimate emails as spam (false positive).

Recall (Sensitivity or True Positive Rate)

- The proportion of **actual positives that are correctly identified**.

- Recall answers the question, "***Of all the actual positive cases, how many did the model correctly identify?***"

$$\text{Recall} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}}$$

When to Use: Use recall **when the cost of false negatives is high.**

Example:

- In disease screening, high recall means most patients with the disease are correctly identified.
- Recall is more important when the cost of **missing a positive case (false negative) is high.**
 - For example, in medical diagnosis (e.g., cancer detection), ***you don't want to miss any potential positive cases (false negatives).***

F1 Score

- The harmonic mean of precision and recall, balancing the two metrics.
- It's harmonic mean of precision & recall.
 - Arithmetic mean is in the centre
 - Harmonic mean is in the lower side
 - It penalises the lower value.**
 - eg. If precision is low, it will penalise it.

$$\text{F1 Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

When to Use:

- Use F1 Score when you need a balance between precision and recall, especially in cases of class imbalance.
- It's useful when you have an imbalanced dataset, where one class is much larger than the other.

When to Use Each Metric:

- **Precision:** Useful when false positives are more costly (e.g., spam detection, fraud detection).
- **Recall:** Important when false negatives are more critical (e.g., cancer detection, disease outbreak).
- **F1 Score:** Ideal when you need to balance both precision and recall, especially in situations where classes are imbalanced.

Python code:

```
from sklearn.metrics import precision_score, recall_score, f1_score

# Example true labels and predicted labels for a binary classification task
y_true = [1, 0, 1, 1, 0, 1, 1, 1, 1, 0]
y_pred = [1, 0, 0, 1, 0, 1, 1, 0, 1, 0]

precision = precision_score(y_true, y_pred)
recall = recall_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred)
```

```
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
```

Output:

Precision: 1.0

Recall: 0.7142857142857143

F1 Score: 0.8333333333333334

We trained 2 models (Logistic regression & Decision Tree) on the heart disease dataset.

```
from sklearn.metrics import recall_score, precision_score, f1_score
```

```
print("For Logistic regression Model")
print("-"*50)
cdf = pd.DataFrame(confusion_matrix(y_test, y_pred1), columns=list(range(0, 2)))
print(cdf)
print("-"*50)
print("Precision - ", precision_score(y_test, y_pred1))
print("Recall - ", recall_score(y_test, y_pred1))
print("F1 score - ", f1_score(y_test, y_pred1))
```

Output:

For Logistic regression Model

```
-----
  0  1
0 82 23
1 10 90
```

```
-----  
Precision - 0.7964601769911505  
Recall - 0.9  
F1 score - 0.8450704225352113
```

```
print("For DT Model")  
print("-"*50)  
cdf = pd.DataFrame(confusion_matrix(y_test,y_pred2),columns=list(range(0,  
2)))  
print(cdf)  
print("-"*50)  
print("Precision - ",precision_score(y_test,y_pred2))  
print("Recall - ",recall_score(y_test,y_pred2))  
print("F1 score - ",f1_score(y_test,y_pred2))
```

Output:

For DT Model

```
-----  
   0   1  
0 101   4  
1   0 100  
-----
```

```
Precision - 0.9615384615384616  
Recall - 1.0  
F1 score - 0.9803921568627451
```

Multi-Class Precision & Recall

- We'll use Iris dataset

```
df = pd.read_csv(r'https://raw.githubusercontent.com/G1Codes/Datasets/refs/heads/main/Iris.csv')
```

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
df['Species'] = encoder.fit_transform(df['Species'])
```

- We have converted Cat → Num

df

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	0
1	2	4.9	3.0	1.4	0.2	0
2	3	4.7	3.2	1.3	0.2	0
3	4	4.6	3.1	1.5	0.2	0
4	5	5.0	3.6	1.4	0.2	0
...
145	146	6.7	3.0	5.2	2.3	2
146	147	6.3	2.5	5.0	1.9	2
147	148	6.5	3.0	5.2	2.0	2
148	149	6.2	3.4	5.4	2.3	2
149	150	5.9	3.0	5.1	1.8	2

TTS:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(df.iloc[:,0:-1],df.iloc[:,-1],test_size
=0.2,random_state=1)
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

clf1 = LogisticRegression()
clf2 = DecisionTreeClassifier()

clf1.fit(X_train,y_train)
clf2.fit(X_train,y_train)

y_pred1 = clf1.predict(X_test)
y_pred2 = clf2.predict(X_test)

from sklearn.metrics import accuracy_score,confusion_matrix
print("Accuracy of Logistic Regression",accuracy_score(y_test,y_pred1))
print("Accuracy of Decision Trees",accuracy_score(y_test,y_pred2))

```

Output:

```

Accuracy of Logistic Regression 1.0
Accuracy of
Decision Trees 0.9666666666666667

```

```

print("Logistic Regression Confusion Matrix\n")
pd.DataFrame(confusion_matrix(y_test,y_pred1),columns=list(range(0,3)))

```

Logistic Regression Confusion Matrix			
	0	1	2
0	11	0	0
1	0	13	0
2	0	0	6

```
print("Decision Tree Confusion Matrix\n")
pd.DataFrame(confusion_matrix(y_test,y_pred2),columns=list(range(0,3)))
```

Decision Tree Confusion Matrix

	0	1	2
0	11	0	0
1	0	12	1
2	0	0	6

```
result = pd.DataFrame()
result['Actual Label'] = y_test
result['Logistic Regression Prediction'] = y_pred1
result['Decision Tree Prediction'] = y_pred2
result.sample(10)
```


	Actual Label	Logistic Regression Prediction	Decision Tree Prediction
31	0	0	0
99	1	1	2
131	2	2	2
125	2	2	2
84	1	1	1
73	1	1	1
40	0	0	0
92	1	1	1
66	1	1	1
90	1	1	1

Now, calculate precision & recall score:

```
from sklearn.metrics import precision_score, recall_score
precision_score(y_test, y_pred1, average=None)
```

Output:

```
array([1., 1., 1.])
```

```
from sklearn.metrics import precision_score, recall_score
precision_score(y_test, y_pred1, average=None)
```

Output:

```
array([1., 1., 1.])
```



If you want average, use `'macro'`

You need to specify an appropriate `average` parameter for multiclass classification.

The `average` parameter can take one of the following values:

- `None` : Returns the recall score for each class.
- `'micro'` : Calculates metrics globally by counting the total true positives, false negatives, and false positives.
- `'macro'` : Calculates metrics for each label and finds their unweighted mean. This does not take label imbalance into account.
 - **Average of all scores**
 - Use macro when the classes are Equal
- `'weighted'` : Calculates metrics for each label and finds their average, weighted by support (the number of true instances for each label). This accounts for label imbalance.
 - You multiply the value with weight of the class
 - Use weighted when classes are imbalanced

To print everything → Use `classification_report`

```
from sklearn.metrics import classification_report
```

```
from sklearn.metrics import classification_report  
print(classification_report(y_test,y_pred1))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	6
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred2))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	11
1	1.00	0.92	0.96	13
2	0.86	1.00	0.92	6
accuracy			0.97	30
macro avg	0.95	0.97	0.96	30
weighted avg	0.97	0.97	0.97	30

- support → How many times 0, 1, 2 occur?