

AdaBoost (Adaptive Boosting)

```
from sklearn.ensemble import AdaBoostClassifier
```

```
from sklearn.ensemble import AdaBoostRegressor
```

- `loss='linear'`

`estimator=None`

`AdaBoostClassifier`

- If `None`, then the base estimator is `DecisionTreeClassifier` initialized with `max_depth=1`.

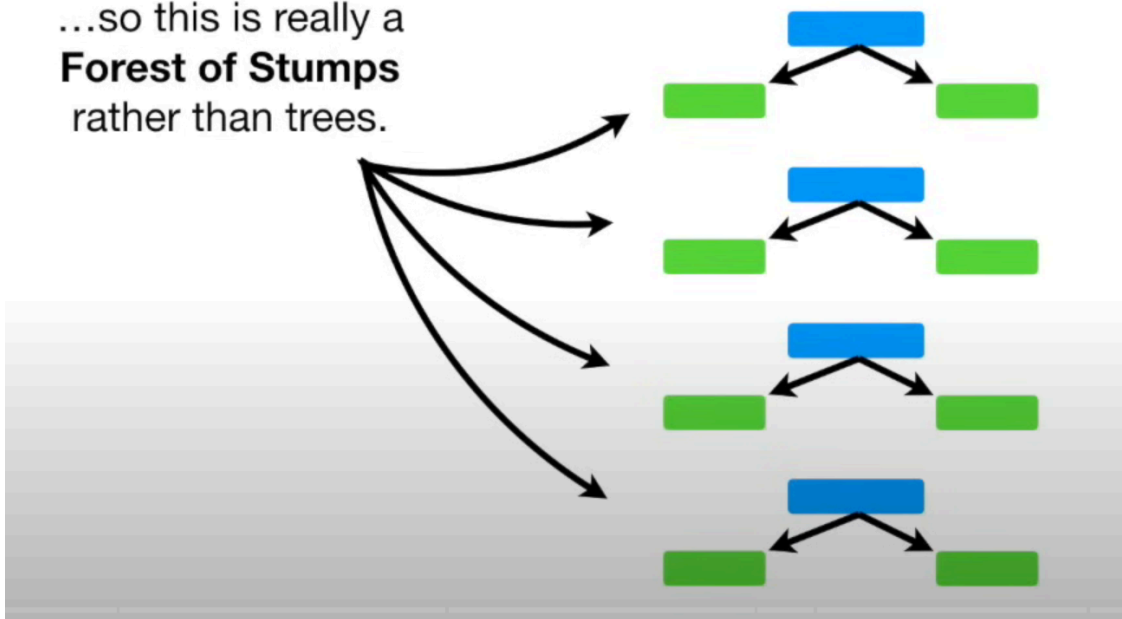
`AdaBoostRegressor`

- If `None`, then the base estimator is `DecisionTreeRegressor` initialized with `max_depth=3`

```
n_estimators=50 , learning_rate=1.0 ,
```

- Combines Decision Trees with a depth of 1, called **Decision Stumps**.
- It works by focusing on the mistakes of previous models and giving more weight to the difficult-to-predict samples.
- **Combines multiple weak classifiers** to form a **strong classifier**.
- It assigns **weights to misclassified samples** to improve future predictions.

...so this is really a
Forest of Stumps
rather than trees.



How Does AdaBoost Work?

1. Train Weak Learner:

- A weak model (e.g., Decision Tree with depth = 1) is trained on the data.

2. Calculate Errors:

- Identify misclassified samples.

3. Adjust Weights:


- Initial weigh $\rightarrow 1/n$
- Increase (**Boost**) the weight of misclassified samples, making them more important for the next weak learner.

4. Train Next Weak Learner:

- The new weak learner focuses more on the misclassified samples.

5. Repeat Process:

- Combine multiple weak learners into a strong learner.



$$\alpha_t = \frac{1}{2} \ln \frac{(1 - TotalError)}{TotalError}$$

- 📌 Formula to calculate the model weight.

Advantages of AdaBoost

1. High Accuracy:

- Often achieves high accuracy by combining multiple weak models.

2. No Need for Parameter Tuning:

- AdaBoost has fewer hyperparameters compared to other algorithms.

3. Handles Both Classification and Regression:

- Can be used for both classification (AdaBoostClassifier) and regression (AdaBoostRegressor).

Disadvantages of AdaBoost

1. **Sensitive to Noisy Data:**

- AdaBoost can overfit if the data contains noise or outliers.

2. **Computationally Expensive:**

- Training multiple models can be computationally expensive.

3. **Requires Weak Models:**

- The performance of AdaBoost depends on the quality of the weak models.

4. **Slower Training:**

- Training multiple models can be computationally expensive.

Why Use AdaBoost?

1. **Improves Accuracy:**

- By combining multiple weak models, AdaBoost can achieve high accuracy.

2. **Handles Complex Data:**

- AdaBoost can capture complex patterns in the data by focusing on difficult samples.

3. **No Need for Deep Trees:**

- Unlike Random Forest, AdaBoost uses very simple models (e.g., Decision Stumps), which are faster to train.

Python code for AdaBoost

Dataset: Telco Customer Churn

gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	TechSupport	StreamingTV	Stres
Female	0	Yes	No	1	No	No phone service	DSL	No	...	No	No	No	
Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes	No	No	
Male	0	No	No	2	Yes	No	DSL	Yes	...	No	No	No	
Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes	Yes	No	
Female	0	No	No	2	Yes	No	Fiber optic	No	...	No	No	No	
...	
Male	0	Yes	Yes	24	Yes	Yes	DSL	Yes	...	Yes	Yes	Yes	
Female	0	Yes	Yes	72	Yes	Yes	Fiber optic	No	...	Yes	No	Yes	
Female	0	Yes	Yes	11	No	No phone service	DSL	Yes	...	No	No	No	
Male	1	Yes	No	4	Yes	Yes	Fiber optic	No	...	No	No	No	
Male	0	No	No	66	Yes	No	Fiber optic	Yes	...	Yes	Yes	Yes	

```

import pandas as pd
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset (replace with your dataset)
# Example: Telco Customer Churn dataset
data = pd.read_csv("https://raw.githubusercontent.com/treselle-systems/customer_churn_analysis/refs/heads/master/WA_Fn-UseC_-Telco-Customer-Churn.csv")

# Preprocess the data (simplified example)
X = data.drop(columns=["Churn"]) # Features
y = data["Churn"] # Target (Churn: Yes/No)

# Convert categorical variables to numerical (if needed)
X = pd.get_dummies(X, drop_first=True)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

```

```
# Create an AdaBoost model with Decision Stumps
ada = AdaBoostClassifier(random_state=42)

# Train the model
ada.fit(X_train, y_train)

# Make predictions
y_pred = ada.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))
```

```
Accuracy: 0.7979176526265973
Classification Report:
              precision    recall  f1-score   support

     No       0.82       0.92       0.87       1539
     Yes       0.68       0.48       0.56        574

 accuracy                0.80       2113
 macro avg       0.75       0.70       0.72       2113
 weighted avg    0.79       0.80       0.79       2113
```

`pd.get_dummies(X, drop_first=True)` is a function in the Pandas library that performs one-hot encoding on categorical variables in a DataFrame x.

- **Categorical Variables:** It identifies categorical columns within the DataFrame x.
- **Creating Dummy Variables:** For each unique category in a categorical column, it creates a new binary column (a "dummy" variable).
- **Binary Representation:** If a row has the specific category in the original column, the corresponding dummy variable gets a value of 1. Otherwise, it gets a value of 0.

`drop_first=True` :

- **Reducing Redundancy:** When you have a categorical variable with n categories, you only need $n-1$ dummy variables to represent it. The n th category can be

inferred when all the other dummy variables are 0.

- **Avoiding Multicollinearity:** In statistical models (like linear regression), including all n dummy variables can lead to multicollinearity, where independent variables are highly correlated. This can cause problems with model estimation.
- **Dropping the First Category:** `drop_first=True` drops the *first* dummy variable that would have been created. This removes the redundancy and avoids multicollinearity.

Let's say you have a DataFrame `X` with a "color" column:

```
color
0    red
1   blue
2  green
3    red
```

When you apply `pd.get_dummies(X, drop_first=True)`, the result would be:

```
color_blue  color_green
0          0          0
1          1          0
2          0          1
3          0          0
```

Is `pd.get_dummies()` Always Necessary?

- **Yes**, if your dataset contains **categorical variables** (text data like "Male" or "Female").
- **No**, if your dataset already contains only numerical data.

Gradient Boosting does not necessarily require `pd.get_dummies` (one-hot encoding)

Select Best Features

```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.5],
    'base_estimator__max_depth': [1, 2, 3]
}

grid_search = GridSearchCV(AdaBoostClassifier(), param_grid, cv=5, scoring='accuracy', n_jobs=-1)
grid_search.fit(X_train, y_train)

print(f"\nBest Parameters: {grid_search.best_params_}")
best_ada = grid_search.best_estimator_
test_accuracy = best_ada.score(X_test, y_test)
print(f"Test Accuracy of Best Model: {test_accuracy:.4f}")

```

 **NOTE:** 🙌 This code hanged the system.

- Took 13+ min to run.

```

Best Parameters: {'learning_rate': 0.5, 'n_estimators': 200}
Test Accuracy of Best Model: 0.8008

```

Python Code for AdaBoost Regression

```

# Step 1: Import necessary libraries
import pandas as pd
from sklearn.ensemble import AdaBoostRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

# Step 2: Load the dataset
from sklearn.datasets import fetch_california_housing
california = fetch_california_housing()

```



```

# Convert the dataset into a pandas DataFrame for better visualization
X = pd.DataFrame(california.data, columns=california.feature_names) # Features
y = pd.Series(california.target) # Target (house prices)

# Step 3: Split the data into training and testing sets
# 80% of the data is used for training, and 20% is used for testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Step 4: Create an AdaBoostRegressor model
# base_estimator = The weak model (Decision Tree with max_depth=3)
# n_estimators = Number of weak models to train
ada = AdaBoostRegressor(n_estimators=100, learning_rate=0.1, random_state=42)

# Step 5: Train the model
ada.fit(X_train, y_train)

# Step 6: Make predictions
y_pred = ada.predict(X_test)

# Step 7: Evaluate the model
# Mean Squared Error (MSE): Lower is better
mse = mean_squared_error(y_test, y_pred)
# R² Score: Closer to 1 is better
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R² Score:", r2)

```

```

Mean Squared Error: 0.568195768061393
R² Score: 0.5663981417281327

```

```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.5]
}

grid_search = GridSearchCV(AdaBoostRegressor(), param_grid, cv=5, scoring
='r2', n_jobs=-1)
grid_search.fit(X_train, y_train)

print(f"\nBest Parameters: {grid_search.best_params_}")
best_ada_reg = grid_search.best_estimator_
r2_best = best_ada_reg.score(X_test, y_test)
print(f"Test R2 with Best Model: {r2_best:.4f}")

```

```

Best Parameters: {'learning_rate': 0.1, 'n_estimators': 50}
Test R2 with Best Model: 0.5684

```



This dataset gives much better result using **GradientBoostingRegressor**

```
GradientBoostingRegressor(n_estimators=200, learning_rate=0.1, max_depth=4, random_state=42)
```

```

Mean Squared Error: 0.2377743906157811
R2 Score: 0.8185494799226947

```

Hyperparameters of AdaBoost

Hyperparameter	Default	What It Does	Effect on Model
base_estimator	None (Defaults to DecisionTreeRegressor(max_depth=3))	Defines the weak learner	More complex base learners

Hyperparameter	Default	What It Does	Effect on Model
		(usually a decision tree).	can lead to overfitting.
n_estimators	50	Number of weak learners (iterations of boosting).	More estimators → Better performance but longer training time.
learning_rate	1	Controls how much each weak model contributes to the final prediction.	Lower values prevent overfitting, but require more estimators.

Recommendation: Use `DecisionTreeRegressor(max_depth=1)`