

# Regularization -Bias Variance Trade-off (VVIMP for Interview)

Asked in ~~almost~~ every interview.

- Only 1% people actually understand this topic
- Can be applied to ALL the ML as well as Deep Learning algorithms.

**Bias (Underfitting)**

**Variance (Overfitting)**



**Regularization Reduces Overfitting (Variance)**

## Bias-Variance Trade-off

Bias and variance are two sources of error in a model:

- **High Bias (Underfitting):** Model is too simple, fails to capture patterns.
- **High Variance (Overfitting):** Model is too complex, captures noise instead of patterns.

💡 **Regularization helps balance bias & variance** by adding a penalty term to the loss function.

## Impact of Regularization

Aspect	Effect on Bias	Effect on Variance
No Regularization	<b>Low Bias:</b> Model fits training data well.	<b>High Variance:</b> Overfits noise in training data.

Aspect	Effect on Bias	Effect on Variance
<b>Moderate Regularization</b>	<b>Balanced:</b> Maintains model flexibility while reducing overfitting.	<b>Reduced:</b> Penalizes complexity, generalizes better.
<b>High Regularization</b>	<b>High Bias:</b> Oversimplifies the model (underfitting).	<b>Low Variance:</b> Model becomes rigid and less sensitive to data fluctuations.

**For adjusting the variance and bias we use:**

1. **Lasso (L1)**
2. **Ridge (L2)**
3. **Elastic Net**

Regularization	Bias	Variance	Effect
No Regularization	Low	High	Overfits
L1 (Lasso)	Medium	Medium	Feature selection + reduced overfitting
L2 (Ridge)	Medium	Lower	Handles collinearity + reduces overfitting
L1 + L2 (Elastic Net)	Medium	Medium	Best of both worlds

## Problem:

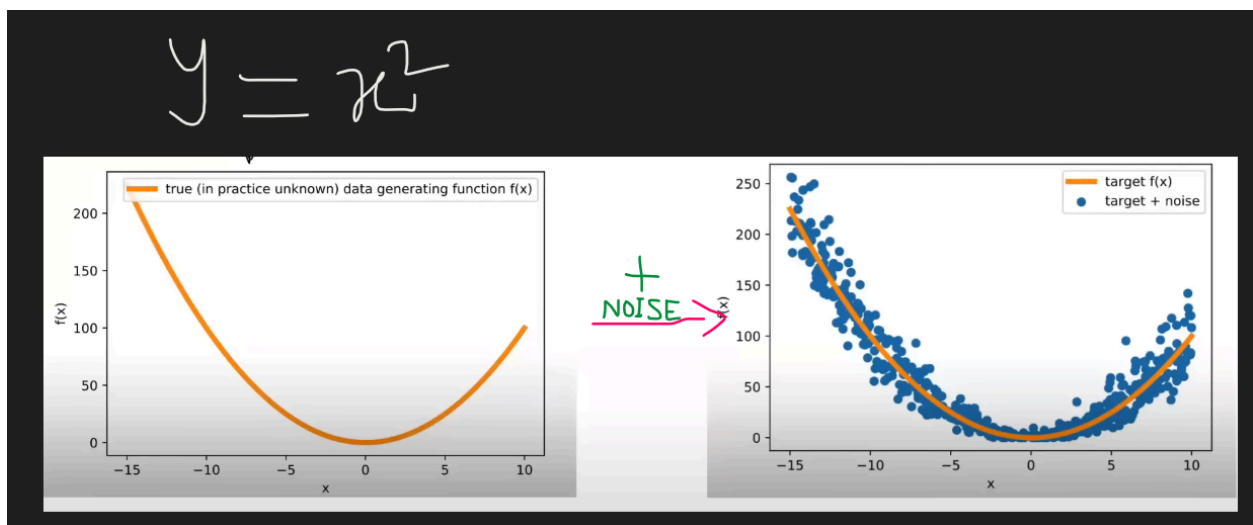
- You have to predict result for the entire population from sample data.
- The WILL BE some error.
- You can't do an accurate prediction due to this error.
- So, we try to find out best estimate.
- If we have data for the entire population, the equation will be  $y = f(x)$
- The prediction is  $\rightarrow f'(x) = \hat{y}$
- $f(x) - f'(x) \rightarrow$  **Reducible Error**
- You can reduce this 🙌 error.

- You cannot reduce the irreducible error.

$$\text{Reducible Error} = \text{Bias}^2 + \text{Variance}$$

## Bias Variance Trade-off

- We'll go reverse:
  - from Population  $\rightarrow$  to Sample
 (In real world, you will not have population data)
- We'll add an error/noise to this




Population data=


$$y = x^2 + \text{error}$$

(-15, 10)

- Draw 3 random samples from the above data 🙌



- We try to fit a linear regression from the sample data:





**BIAS:** The inability of a machine learning model to fit the training data

### Underfitting

- Above graph is **HIGH BIAS**.

High Bias = Underfitting

- As the bias decreases, the data starts fitting

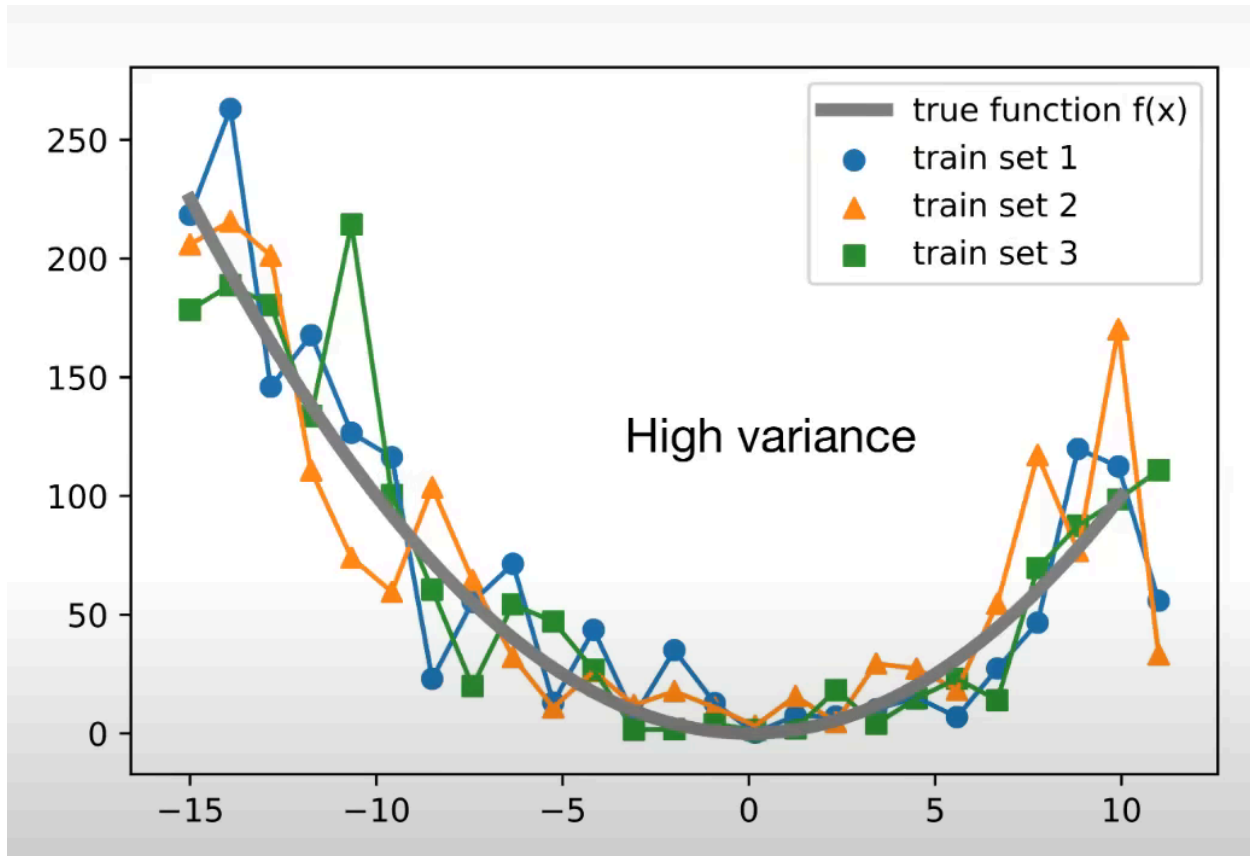


**Variance:** Change in ML model when data is changed.

### Overfitting 🙅

- Above graph is **LOW VARIANCE**.
- **Variance** in machine learning refers to the model's **sensitivity to small changes in the training data**.
- A model with high variance overfits, meaning it **captures noise and random fluctuations** in the training data, leading to **poor generalization** to new data.

**Now, apply polynomial regression to the above data.**



- The models are Low Bias (less Underfitting)
  - The training data is fitting very well.
- But the results of all three models it's varying from each other.
  - Therefore it's a **high variance** model.



**High variance** is closely related to **overfitting**.




**Underfitting** is closely related to **high bias**.

## Ideal Situation:

### Low variance-Low Bias

**Meaning** → Your data is fitting test data properly. And when you get another training data, the model doesn't drastically change.

- **Problem** → Bias and Variance are inversely proportional to each other.



- As you increase the complexity, bias decreases
- BUT variance starts to increase.
- **AIM : To find the middle point**

## Expected Value and Variance

- Expected value represents the average outcome of a random variable over a large number of trials or experiments.

- We roll a die 1 Lac times
  - Mean will be  $\rightarrow 3.5$
  - This is EXPECTED VALUE

**Expected Value  $E[X]$  = Population Mean**

**$\text{Var}(X)$  = Variance of Population**

$$\begin{aligned}
 \text{Var}(X) &= E[(X - E[X])^2] \\
 &= E[X^2 - 2XE[X] + E[X]^2] \\
 &= E[X^2] - 2E[X]E[X] + E[X]^2 \\
 &= E[X^2] - 2E[X]^2 + E[X]^2 \\
 &= E[X^2] - E[X]^2
 \end{aligned}$$

## Bias and Variance Mathematically?

**Bias:**

$$\text{Bias}(f'(x)) = E[f'(x)] - \underline{f'(x)}$$



$f(x) \rightarrow$  Population mean (Expected value of the population)

$f'(x) \rightarrow$  Sample mean

- If difference between them is zero  $\rightarrow$  Our model is unbiased



**If we draw 100 samples & find out the mean  $\rightarrow$  It will be close to the population mean.**

## Variance:

- Variance refers to the amount by which the prediction of our model will change if we used a different training data set.
- In other words, it measures how much the predictions for a given point vary between different realizations of the model.

$$\text{Var}(f'(x)) = E[(f'(x) - E[f'(x)])^2]$$

- **If this is high  $\rightarrow$  Upon changing the data, the accuracy, R2 score, etc will change a lot.**
- Because the model is **OVERFITTING**

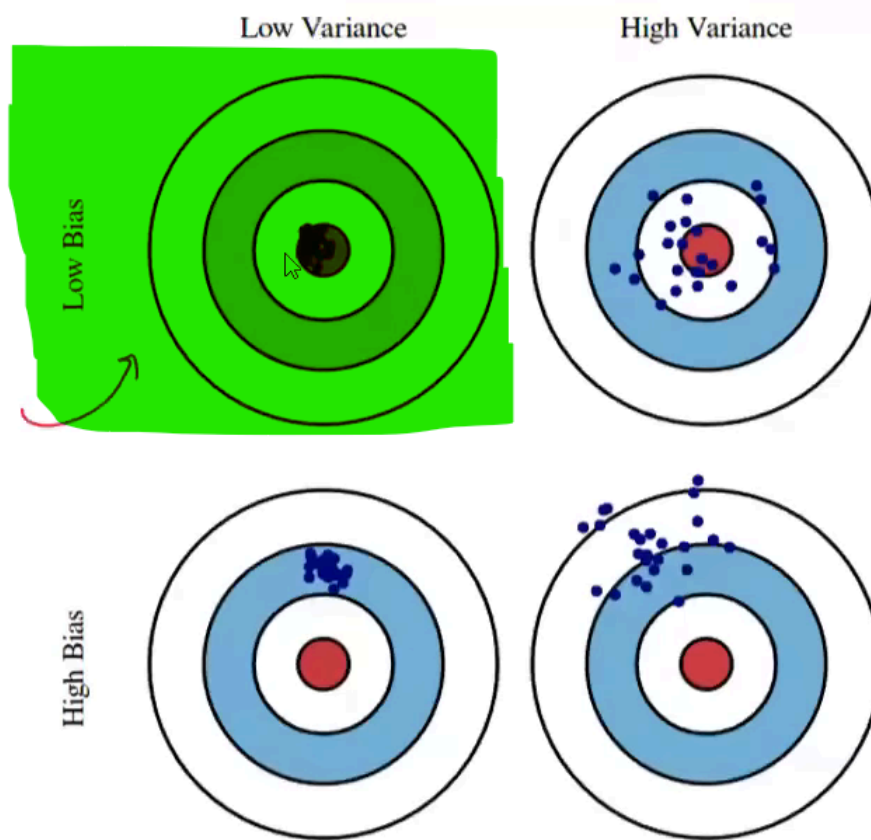


Fig. 1 Graphical illustration of bias and variance.




Fig. 1 Graphical illustration of bias and variance.

## Bias Variance Decomposition (VIMP)

- It divides the loss (eg. MSE) into 3 parts:
  1. Bias
  2. Variance
  3. Irreducible Error

$$\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$



**VIMP NOTE:** Here, we're DECOMPOSING. First we get the loss. Then we split it into these 3.

- **Irreducible Error ( $\epsilon$ )** → Moving target/Noise

**Bias<sup>2</sup> will cancel out the negative values.**

- **Bias + Variance** → Reducible Errors
- **High Bias:** Simplified model → Underfitting → Low accuracy.
- **High Variance:** Complex model → Overfitting → Poor generalization.

## Code Example

```
from mlxtend.evaluate import bias_variance_decomp
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from mlxtend.data import boston_housing_data
from sklearn.model_selection import train_test_split
```

```
X, y = boston_housing_data()
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.3,
                                                    random_state=123,
                                                    shuffle=True)
```

```
lr = LinearRegression()

avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
    lr, X_train, y_train, X_test, y_test,
    loss='mse',
    random_seed=123)

print('Average expected loss: %.3f' % avg_expected_loss)
```

```
print('Average bias: %.3f' % avg_bias)
print('Average variance: %.3f' % avg_var)
```

### Output:

Average **expected loss**: 29.891    **#(MSE)**

Average **bias**: 28.609                      **#Bias<sup>2</sup>**

Average **variance**: 1.282

**bias\_variance\_decomp(...)** :

- This function is used to compute the **bias**, **variance**, and **expected loss** of the model on a given dataset.
  - **avg\_expected\_loss** : The average error on the test set (MSE).
  - **avg\_bias** : The average bias of the model (how far off the model's predictions are from the true values).
  - **avg\_var** : The variance of the model's predictions (how much the predictions fluctuate across different data subsets).

**Total Error = Bias<sup>2</sup> + Variance + Irreducible Error**

- Linear Regression gives **high bias** (predictions are away from actual value) &
- **Low variance** (It's precise when ran multiple times)
- **DecisionTreeRegressor** is opposite of linear regression

**Let's apply **DecisionTreeRegressor** on same dataset:**

```
dt = DecisionTreeRegressor(random_state=123)

avg_expected_loss, avg_bias, avg_var = bias_variance_decomp(
```

```
dt, X_train, y_train, X_test, y_test,  
loss='mse',  
random_seed=123)  
  
print('Average expected loss: %.3f' % avg_expected_loss)  
print('Average bias: %.3f' % avg_bias)  
print('Average variance: %.3f' % avg_var)
```

### **Output:**

Average expected loss: 31.536

Average bias: 14.096

Average variance: 17.440

- Here, bias has reduced but variance is increased as compared to linear regression

## When to use Regularization?

### 1. Prevent Overfitting

- Use when your model performs well on training data but poorly on validation/test data.

### 2. High Dimensionality (Many Features, Few Samples)

- Use L1/L2 to reduce model complexity and avoid overfitting (e.g., text/gene data).

### 3. Multicollinearity (Correlated Features)

- Use **Ridge (L2)** to stabilize coefficients and distribute weights among correlated features.

### 4. Feature Selection

- Use **Lasso (L1)** to shrink irrelevant features' coefficients to zero, retaining only important ones.

### 5. Improve Interpretability

- Simplify models by reducing feature count (L1) or shrinking coefficients (L2).

## 6. **Boost Model Performance**

- Apply regularization to enhance out-of-sample performance, even if overfitting isn't evident.