

SVM

Used for:

- **Classification** → **Binary + Multiclass**
- **Regression**
- **Image processing**
- Works for both Linear + Non-linear

Drawback:

- Difficult to understand



Gives good results in almost any situation.

Goal: to find a **hyperplane** that best separates the data into classes with the maximum possible margin.

How Does SVM Work?

SVM works by:

- ✓ **Finding a hyperplane** that separates classes
- ✓ **Maximizing the margin** (distance between the hyperplane and the nearest data points)
- ✓ **Using support vectors** (points that determine the boundary)

Key Concepts

Hyperplane:

- A decision boundary that separates data points of different classes.
- In 2D, it's a line; in 3D, it's a plane; in higher dimensions, it's a hyperplane.

Support Vectors:

- The **closest data points** to the hyperplane.
- These **determine the position and orientation** of the hyperplane.
- The model **only depends on these points**.

Margin:

- The distance between the hyperplane and the nearest data points (support vectors).
- SVM aims to maximize this margin.

Types of SVM

1 Linear SVM → Used when data is **linearly separable**.

- Assumes data is linearly separable.
- Finds a straight line (or hyperplane) to separate classes.

2 Non-Linear SVM (Kernel SVM) → Used when data is **not linearly separable**.

- Uses **kernel functions** to transform data into a higher-dimensional space where it becomes linearly separable.
- Common kernels: Polynomial, Radial Basis Function (RBF), Sigmoid.

SVM Parameters

1. **C (Regularization Parameter):**

- **C** controls the trade-off between achieving a low error on the training data and maintaining a large margin.
 - **High C:** SVM tries to classify all training points correctly, which may lead to overfitting (narrow margin).
 - **Low C:** SVM allows some misclassifications but aims for a wider margin, which might lead to underfitting.

2. Kernel:

- Specifies the type of kernel used (e.g., **linear**, **RBF**, **polynomial**, etc.).

3. Gamma (for non-linear kernels like RBF):

- **Gamma** defines how far the influence of a single training example reaches. A low gamma means a far-reaching influence, while a high gamma means a closer, more local influence.
 - **Low gamma:** The decision boundary is smoother.
 - **High gamma:** The decision boundary is more flexible and could overfit the data.

4. Degree (for polynomial kernel):

- Specifies the degree of the polynomial kernel function (used only when kernel = polynomial).

Classification

1. **Hard Margin SVM** → assumes **perfect separation** (works only for clean datasets).
2. **Soft Margin SVM** → allows **some misclassification** using a penalty term **C**.

Support Vector Classifier

C (Regularization Parameter):

- **High C** → Tries to **classify every point correctly** (low bias, high variance).
- **Low C** → Allows **some misclassification** (high bias, low variance).


3. SVM Kernels → Non-linear


4. SVM for multi-class

5. SVR → Regression

Hard Margin SVM (Maximal Margin Classifier)


ANALYTIX LABS






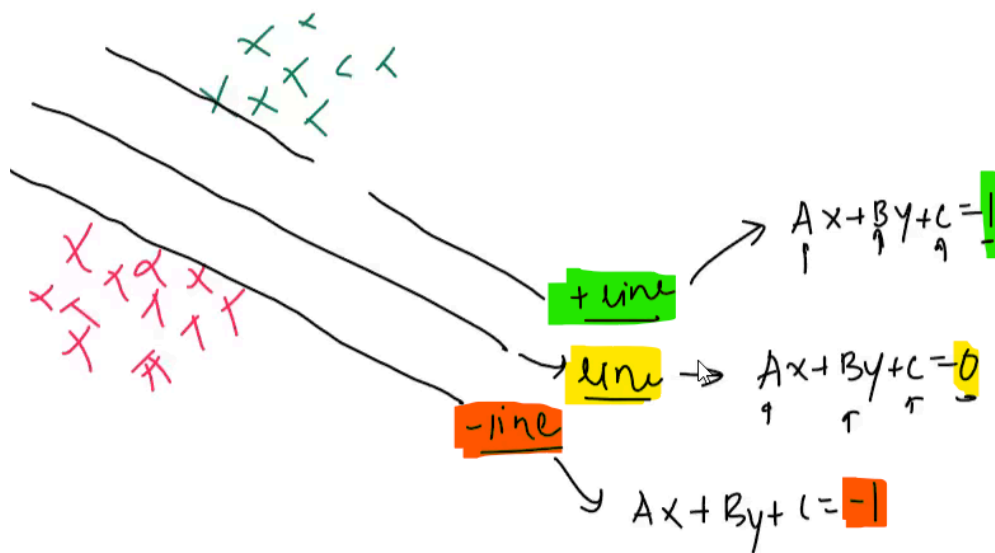
Support Vectors:

- Vectors which lie on the outer line





- In **Hard Margin SVM**, we do not want any points inside the support vector lines.



- 👉 We can **multiply** the equation by a **number more than 1** to **decrease** the margin.
- Multiply with **number less than 1** to **increase the distance**.

$$Ax + By + c \geq 1$$

$$Ax + By + c \leq -1$$

- We want to maximize the distance between support vectors which satisfies the below equation:

$$\left\{ \begin{array}{l} \text{if } y_i = 1 \quad Ax_{i1} + Bx_{i2} + C \geq 1 \\ y_i = -1 \quad Ax_{i1} + Bx_{i2} + C \leq -1 \end{array} \right.$$

- We can write the above equation as:

$$y_i (Ax_{i1} + Bx_{i2} + C) \geq 1$$

$\begin{matrix} +ve \\ -ve \end{matrix}$

Loss Function:

You maximize this 📌

$$\operatorname{argmax}_{A, B, C} \frac{2}{\sqrt{A^2 + B^2}} \quad \text{given} \left\{ y_i (Ax_{i1} + Bx_{i2} + C) \geq 1 \right\}$$

	x_{2i}	y_i
	\uparrow	\uparrow
x_1	x_2	y
x_{11}	x_{21}	y_1
x_{12}	x_{22}	y_2
x_{13}	x_{23}	y_3

- This is constrained optimization problem because we have to solve 2 things at a time.
- We solve this with quadratic programming (high level math we won't get into)



You never use hard margin SVM because real world datasets are not perfectly separable.

Drawbacks of Hard Margin SVM:

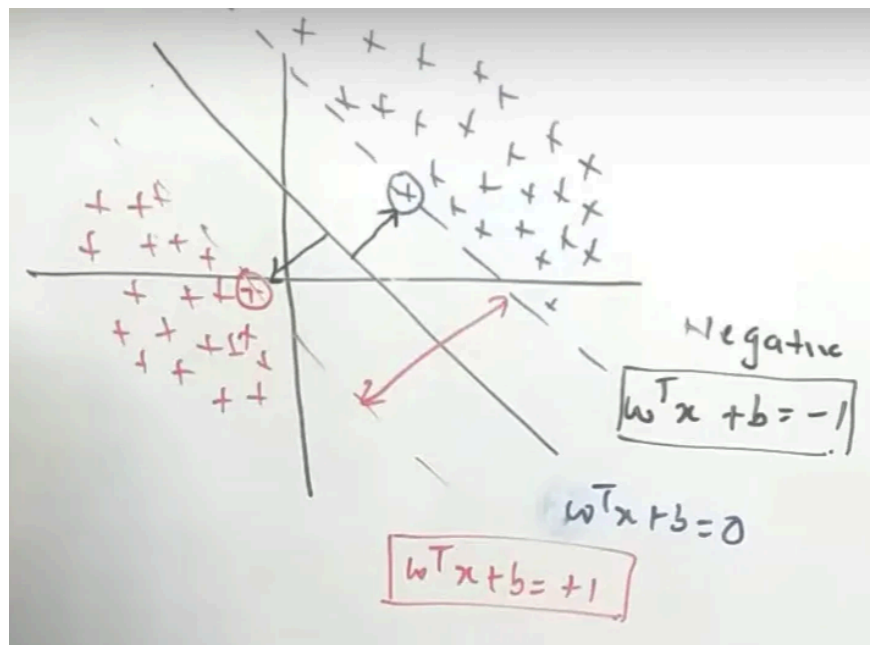
- Sensitive to outliers.
- **Limited to Linearly Separable Data:**
 - Hard margin SVMs can only be used when the data is perfectly linearly separable. In real-world scenarios, data is often complex and non-linearly separable, rendering hard margin SVMs impractical.
- Lack of Flexibility.

Soft Margin vs. Hard Margin

Feature	Hard Margin SVM	Soft Margin SVM
Misclassification Allowed?	✗ No	✓ Yes
Works for Noisy Data?	✗ No	✓ Yes
Overfitting Risk	✓ High	✗ Lower
Used in Real-World?	✗ Rarely	✓ Yes

Soft Margin SVM (SVC)

- **Soft Margin SVM** allows **some misclassification** using a penalty term **C**.
- Here, we **soften the constraint**.



Slack Variable


- Slack variables (often denoted as ξ or ζ) are introduced to relax the strict requirement of perfect separation.
- They allow some data points to violate the margin or even be misclassified.



- Each data point is associated with a slack variable, representing the degree to which it violates the margin.

How Slack Variables Work:

- When a data point is correctly classified and lies outside the margin, its slack variable is zero.
- When a data point lies within the margin but is correctly classified, its slack variable is greater than zero but less than one.
- When a data point is misclassified, its slack variable is greater than one.

This balance is controlled by a regularization parameter (often denoted as C)





- $\xi(C/\text{Zeta})$ is Misclassification score.
 - Also called as Hinge loss in ML.

Mathematical Formulation of Soft Margin SVM:

$$\min_{w,b,\xi} \frac{1}{2} ||w||^2 + C \sum_{i=1}^N \xi_i$$

Where:

- w : weight vector perpendicular to the hyperplane.
- b : bias term of the hyperplane.
- ξ_i : slack variable for the i -th data point.
- C : regularization parameter controlling the penalty for misclassification .

$$y_i (Ax_{1i} + Bx_{2i} + C) \geq 1 - \xi_i$$

such that

$$\xi_i \geq 0$$

$$\boxed{X_1(Ax_{1i} + Bx_{2i} + C) \geq 1 - \xi_i} \rightarrow \begin{array}{l} \text{it is allowing} \\ \text{all the points} \end{array}$$

\hookrightarrow this is no more a constraint

\hookrightarrow All true condition

- This allowed all the points.
- It's no more a constraint.
- So, we add this 📌 thing to the above formula:

$$\sum_{i=1}^N \xi_i$$

Objective Function:

$$\min_{w,b,\xi} \frac{1}{2} ||w||^2 + C \sum_{i=1}^N \xi_i$$

Where:

- w : weight vector perpendicular to the hyperplane.
- b : bias term of the hyperplane.
- ξ_i : slack variable for the i -th data point.
- C : regularization parameter controlling the penalty for misclassification.

Subject to Constraints:

$$y_i(w \cdot x_i + b) \geq 1 - \xi_i \quad \text{for } i = 1, 2, \dots, N$$

$$\xi_i \geq 0 \quad \text{for } i = 1, 2, \dots, N$$

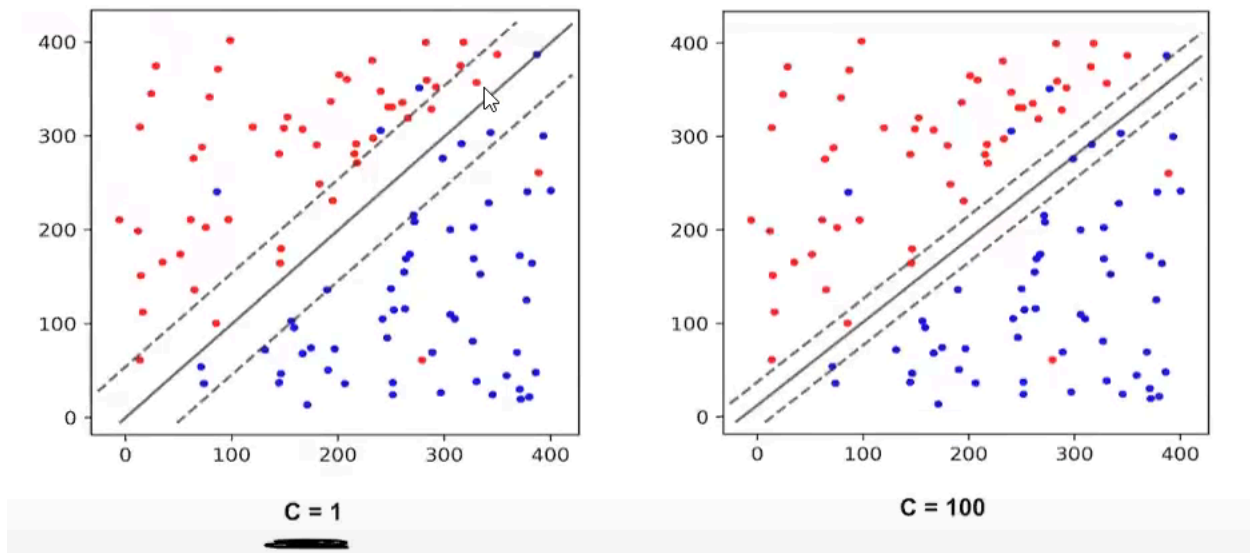
Where:

- $y_i \in \{-1, 1\}$: class label of the i -th data point.
- x_i : feature vector of the i -th data point.
- ξ_i : slack variable.

ξ_i = **Slack variable (allows misclassification)**

C = **Regularization parameter (controls trade-off)**

- **High C** → Tries to **classify every point correctly (low bias, high variance)**.
- **Low C** → Allows **some misclassification (high bias, low variance)**.



Python Code

```
# Import necessary libraries
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# Load the Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Only take two classes (binary classification for simplicity)
X = X[y != 2]
y = y[y != 2]

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Initialize the SVM classifier with a linear kernel (soft margin)
svm_model = SVC(kernel='linear', C=1) # C is the regularization parameter


# Train the model
svm_model.fit(X_train, y_train)

# Make predictions
y_pred = svm_model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy * 100:.2f}%")
```

Accuracy: 100.00%

Relationship with Logistic Regression



C (Regularization Parameter) & ξ (Slack Variable)

C (Regularization Parameter)

C is a **hyperparameter** that you choose before training the model

- Controls the trade-off between **maximizing the margin** and **minimizing misclassification**.
- **High C** → **Less misclassification**, smaller margin, risk of **overfitting**.
- **Low C** → **More misclassification allowed**, larger margin, better **generalization**.

ξ (Slack Variable)

ξ are **variables** created for each data point during the model training process

- Measures how much a point **violates** the margin.
- If $\xi = 0$, the point is correctly classified and outside the margin.
- If $0 < \xi \leq 1$, the point is inside the margin but correctly classified.
- If $\xi > 1$, the point is misclassified.
- ξ is not something you choose; it's computed during training as part of the optimization process.